

Exercise 2:

Aim:

To understand and implement real-time data communication using Apache Kafka by building a Python producer that sends messages and a Python consumer that receives them via a Kafka topic. This exercise demonstrates core messaging concepts such as publish-subscribe, message queuing, brokers, and topic-based communication.

Procedure:

Phase 1:

- Set Up Kafka Broker Using Docker Compose
- Create a folder named **Kafka_Lab** on your Desktop.
- Inside the folder, create a **docker-compose.yml** file defining Kafka and Zookeeper services.
- Run **docker-compose up -d** to start Kafka and Zookeeper containers.
- Verify the setup with **docker ps** to ensure both containers are running.

Phase 2:

- Create a Kafka Topic
- Enter the Kafka container using: **docker exec -it kafka /bin/bash**
- create a topic named user-events:
 - **kafka-topics --create --topic user-events --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1**
- **Exit**

Phase 3:

- Set Up Python Environment
- Create and activate a virtual environment:
 - **python -m venv kafka_env**
 - **kafka_env\Scripts\activate** # Windows
- Install the Kafka Python client:
 - **pip install kafka-python**

Phase 4:

- Create producer.py to simulate and send JSON-encoded user data to Kafka.
- Create consumer.py to read and display the JSON messages from the Kafka topic.

Phase 5:

- Run and Test the System
- In one terminal, run the consumer:
 - **python consumer.py**
- In another terminal, run the producer:
 - **python producer.py**
- Observe how messages sent by the producer are consumed and printed in real-time by the consumer.

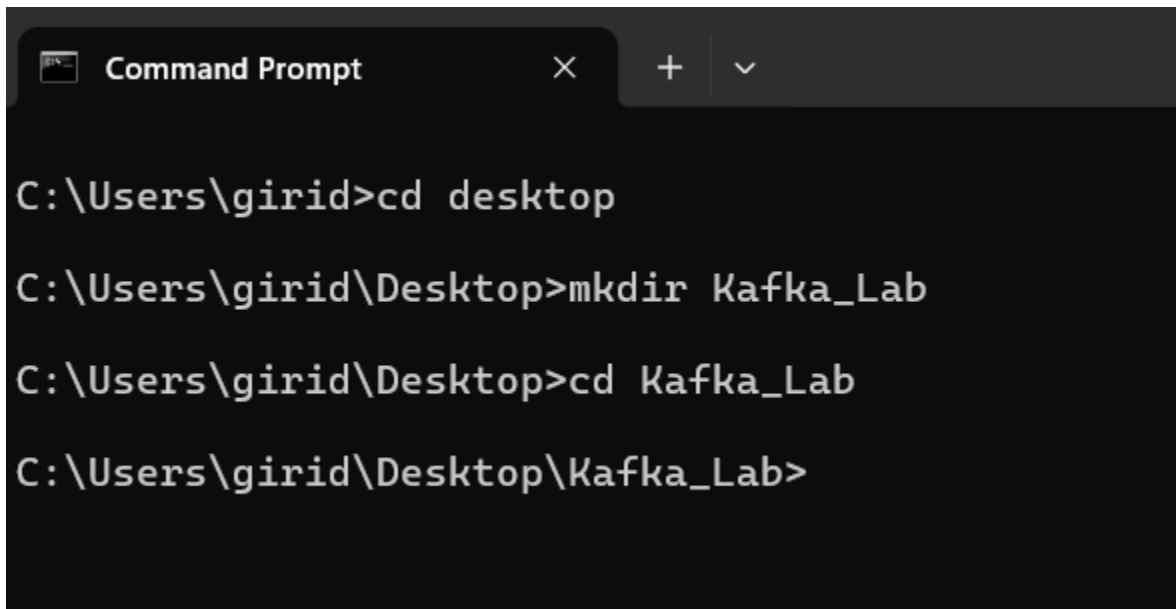
Phase 6:

- To stop and remove Kafka and Zookeeper containers, run:
 - **docker-compose down**

Output:

Phase 1:

- Creating folder.



```
Command Prompt
C:\Users\girid>cd desktop
C:\Users\girid\Desktop>mkdir Kafka_Lab
C:\Users\girid\Desktop>cd Kafka_Lab
C:\Users\girid\Desktop\Kafka_Lab>
```

- Creating docker-compose.yml file inside Kafka_Lab:

```
# version: '3.8'

services:
  zookeeper:
    image: confluentinc/cp-zookeeper:7.4.0
    container_name: zookeeper
    ports:
      - "2181:2181"
    environment:
      ZOOKEEPER_CLIENT_PORT: 2181
      ZOOKEEPER_TICK_TIME: 2000

  kafka:
    image: confluentinc/cp-kafka:7.4.0
    container_name: kafka
    ports:
      - "9092:9092"
    depends_on:
      - zookeeper
    environment:
      KAFKA_BROKER_ID: 1
      KAFKA_ZOOKEEPER_CONNECT: zookeeper:2181
      KAFKA_ADVERTISED_LISTENERS: PLAINTEXT://localhost:9092
      KAFKA_OFFSETS_TOPIC_REPLICATION_FACTOR: 1
      KAFKA_LISTENERS: PLAINTEXT://0.0.0.0:9092
```

- compose up and docker ps

```
C:\Users\girid\Desktop\Kafka_Lab>docker-compose up -d
time="2025-07-23T16:19:26+05:30" level=warning msg="C:\\Users\\girid\\Desktop\\Kafka_Lab\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion"
[+] Running 3/3
 ✓ Network kafka_lab_default Created 0.0s
 ✓ Container zookeeper Started 0.5s
 ✓ Container kafka Started 0.5s

C:\Users\girid\Desktop\Kafka_Lab>docker ps
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS                                NAMES
b3f973f78b96   confluentinc/cp-kafka:7.4.0        "/etc/confluent/dock_  7 seconds ago Up 7 seconds  0.0.0.0:9092->9092/tcp, [::]:9092->9092/tcp   kafka
336e5528e494   confluentinc/cp-zookeeper:7.4.0    "/etc/confluent/dock_  8 seconds ago Up 7 seconds  0.0.0.0:2181->2181/tcp, [::]:2181->2181/tcp   zooke

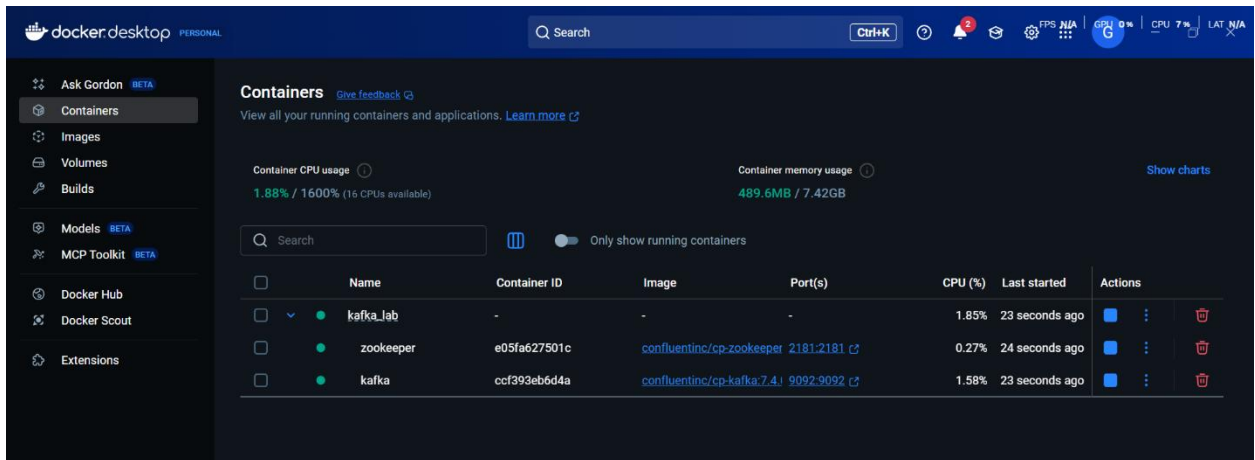
C:\Users\girid\Desktop\Kafka_Lab>docker exec -it kafka /bin/bash
[appuser@b3f973f78b96 ~]$
```

Phase 2:

- Kafka container

```
C:\Users\girid\Desktop\Kafka_Lab>docker exec -it kafka /bin/bash
[appuser@b3f973f78b96 ~]$
```

```
[appuser@b3f973f78b96 ~]$ kafka-topics --create --topic user-events --bootstrap-server localhost:9092 --partitions 1 --replication-factor 1
Created topic user-events.
[appuser@b3f973f78b96 ~]$
```



Phase 3:

- Creating python venv

```
C:\Users\girid\Desktop\Kafka_Lab>python -m venv kafka_env
C:\Users\girid\Desktop\Kafka_Lab>kafka_env\Scripts\activate
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>
```

Installing kafka-python:

```
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>pip install kafka-python
Collecting kafka-python
  Downloading kafka_python-2.2.15-py2.py3-none-any.whl.metadata (10.0 kB)
  Downloading kafka_python-2.2.15-py2.py3-none-any.whl (309 kB)
Installing collected packages: kafka-python
Successfully installed kafka-python-2.2.15

[notice] A new release of pip is available: 24.2 -> 25.1.1
[notice] To update, run: python.exe -m pip install --upgrade pip

(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>
```

Phase 4:

producer.py:

```
import time
import json
import random
from kafka import KafkaProducer
producer = KafkaProducer(
    bootstrap_servers=['localhost:9092'],
    # Encode all values as JSON
    value_serializer=lambda v: json.dumps(v).encode('utf-8')
)

print("Producer started. Press Ctrl+C to stop.")

user_id_counter = 1

try:
    while True:
        user_data = {
            'user_id': f'user_{user_id_counter}',
            'event_type': 'account_creation',
            'username': f'user{user_id_counter}',
            'timestamp': time.time(),
            'country_code': random.choice(['US', 'CA', 'GB', 'AU', 'DE'])
        }
        producer.send('user-events', value=user_data)
        print(f"Sent event for user_id: {user_data['user_id']}")
        user_id_counter += 1
        time.sleep(random.uniform(1, 3))
```

```

except KeyboardInterrupt:
    print("Stopping producer.")
finally:
    producer.flush()
    producer.close()
    print("Producer closed.")

```

consumer.py:

```

import json
from kafka import KafkaConsumer
consumer = KafkaConsumer(
    'user-events', # The topic to subscribe to
    bootstrap_servers=['localhost:9092'],
    auto_offset_reset='earliest', # Start from the earliest message
    value_deserializer=lambda m: json.loads(m.decode('utf-8')) # Decode JSON
    messages
)

print("Consumer started. Waiting for messages... (Press Ctrl+C to stop)")

try:
    for message in consumer:
        event_data = message.value # Deserialized JSON data
        print("\n-----")
        print("Received new event:")
        print(f" Topic: {message.topic}")
        print(f" Partition: {message.partition}, Offset: {message.offset}")
        print(f" User ID: {event_data.get('user_id')}")
        print(f" Event Type: {event_data.get('event_type')}")
        print(f" Username: {event_data.get('username')}")
        print("-----")
except KeyboardInterrupt:
    print("\nStopping consumer.")
finally:
    consumer.close()
    print("Consumer closed.")

```

Phase 5:

- **producer.py:**

```
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>python producer.py
Producer started. Press Ctrl+C to stop.
Sent event for user_id: user_1
Sent event for user_id: user_2
Sent event for user_id: user_3
Sent event for user_id: user_4
Sent event for user_id: user_5
Sent event for user_id: user_6
Sent event for user_id: user_7
|
```

- **consumer.py**

```
C:\Users\girid>cd desktop
C:\Users\girid\Desktop>cd Kafka_Lab
C:\Users\girid\Desktop\Kafka_Lab>kafka_env\Scripts\activate
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>python consumer.py
Consumer started. Waiting for messages... (Press Ctrl+C to stop)

-----
Received new event:
Topic: user-events
Partition: 0, Offset: 0
User ID: user_1
Event Type: account_creation
Username: user1
-----

-----
Received new event:
Topic: user-events
Partition: 0, Offset: 1
User ID: user_2
Event Type: account_creation
Username: user2
-----

-----
Received new event:
Topic: user-events
Partition: 0, Offset: 2
User ID: user_3
Event Type: account_creation
Username: user3
-----
```

Phase 6:

- Removing kafka containers

```
Stopping producer.  
Producer closed.  
  
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>docker-compose down  
time="2025-07-23T17:46:05+05:30" level=warning msg="C:\\Users\\girid\\Desktop\\Kafka_Lab\\docker-compose.yml: the attribute 'version' is obsolete, it will be ignored, please remove it to avoid potential confusion"  
[+] Running 3/3  
  ✓ Container kafka          Removed      1.0s  
  ✓ Container zookeeper      Removed      0.7s  
  ✓ Network kafka_lab_default Removed      0.3s  
  
(kafka_env) C:\Users\girid\Desktop\Kafka_Lab>
```